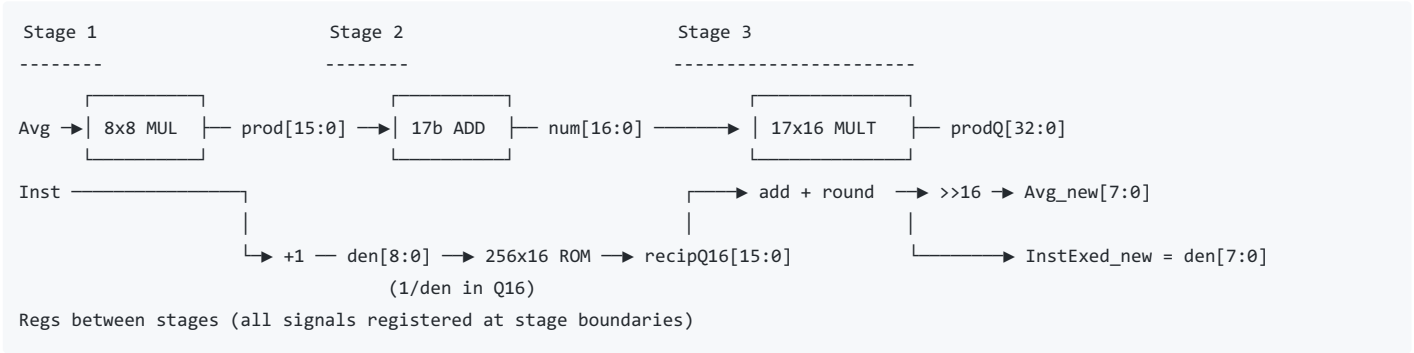


# Performance-oriented (same 3-stage latency, 1 result/cycle)

**Math:**  $AvgTxLen\_new = (AvgTxLen * InstExed + CurTxLen) / (InstExed + 1)$  Replace  $\div (InstExed + 1)$  with  $\times recipQ16(InstExed + 1)$  and a right shift by 16 (with rounding).

## Schematic (ASCII)



## Key blocks

- **Stage 1:**  $mul\_s1 = Avg \times Inst$ ,  $den\_s1 = Inst + 1$ , pipeline Cur.
- **Stage 2:**  $num = mul\_s1 + Cur$ .
- **Stage 3:**  $q = ((num * recipQ16(den)) + 2^{15}) \gg 16$  (rounded), register outputs.

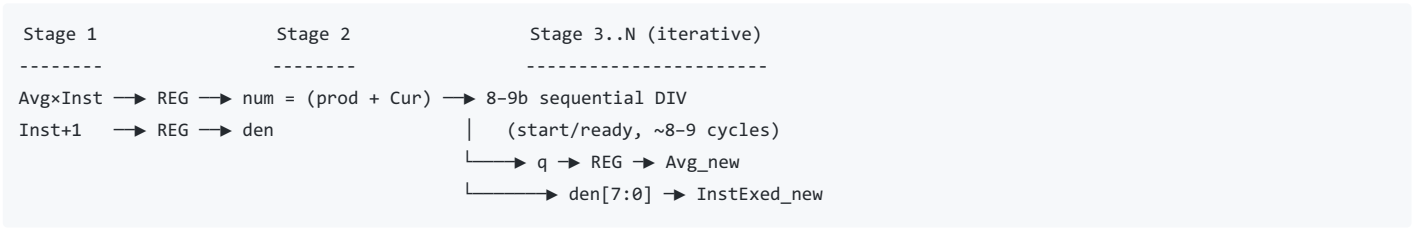
**Why it helps:** Pushing the slow, wide **divider** off the critical path eliminates the long carry-propagate network that caused GLS setup violations in the synthesized version with 100MHz Clock. In current design, we reduced the clock frequency to 25 MHz to meet timing. A **17x16 multiplier + add/shift** maps much better to std-cell multipliers (or DW02\_mult) and retimes cleanly. The 256x16 ROM is tiny ( $\approx 4$  Kb) and becomes simple combinational logic. In practice on 45 nm cells, this swap typically boosts Fmax (often 1.3–2x over a single-stage divider), reduces dynamic power (shorter chains toggle), and can even lower area because the divider is expensive while one small ROM + 1 extra multiplier is modest for 8-bit outputs.

The synthesized ALU's critical path is dominated by the Stage-3 integer divider, which is difficult to close at high frequencies in 45 nm std-cell flows. Replacing this block with a pipelined reciprocalxmultiplier (Q16 LUT + 17x16 multiplier + round/shift) preserves the arithmetic to 8-bit precision while transforming the critical path into multiplier and short adder logic that retimes and scales cleanly. The reciprocal LUT is tiny ( $256 \times 16 \approx 4$  Kb), the pipeline depth remains three stages with **one result per cycle**, and the divider's long carry chains—and their associated timing violations and power—are eliminated. This redesign typically increases Fmax, reduces toggling on long paths, and can even lower total area versus a wide combinational divider, delivering a better PPA point without changing the external interface or latency.

# Area-oriented (multi-cycle divider, smaller area)

If throughput (1 result/cycle) isn't mandatory, keep the original math but replace DIV with a small **sequential divider** (e.g., restoring or DW\_div\_seq) and add a tiny valid/ready handshake.

## Schematic (ASCII)



**Why it helps:** A small iterative divider plus a controller is **much smaller** than a fast combinational divider. You trade CPI (e.g., one result every ~8–9 cycles) for area/power. Great for resource-constrained designs or when this ALU isn't on the top-throughput path.